

AirBnb Ratings Classification

Matthew Zhang

Objective:

Often times AirBnb success is determined by customer experience. In this case, I will be a company analyst looking to help AirBnb hosts enhance overall customer satisfaction by classifying listings into three categories: Subpar(0), Good(1) and Best(2); further, determining which aspects of properties can be improved. Through modeling, I hope to gain valuable insights in order to form strategy and help generate company success.

To illustrate which characteristics of listings are important, I hope to perform feature level analysis to gain further insights into the models and how they behave.

Data Cleaning

```
In [1]: #Import necessary packages
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, mean_squared_error, plot_confusion_matrix
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
from sklearn.preprocessing import StandardScaler

import warnings
warnings.filterwarnings('ignore')
```

```
In [9]: #Import relevant dataset
df = pd.read_csv('data/listings (1).csv.gz')
rev = pd.read_csv('data/reviews.csv.gz')
```

```
In [11]: pd.set_option('display.max_columns', None)
df.head(2)
```

Out[11]:

	id	listing_url	scrape_id	last_scraped	name	description
0	109	https://www.airbnb.com/rooms/109	20201103044428	2020-11-03	Amazing bright elegant condo park front *UPGRA...	*** Unit upgraded with new bamboo flooring, br...
1	2708	https://www.airbnb.com/rooms/2708	20201103044428	2020-11-04	Beautiful Furnish Mirrored Mini-Suite w/ Firep...	CDC Airbnb Standard Steam, Sanitized, Disinf...

```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31471 entries, 0 to 31470
Data columns (total 74 columns):
id                                     31471 non-null int64
listing_url                           31471 non-null object
scrape_id                             31471 non-null int64
last_scraped                          31471 non-null object
name                                  31469 non-null object
description                            30400 non-null object
neighborhood_overview                 20323 non-null object
picture_url                           31471 non-null object
host_id                               31471 non-null int64
host_url                              31471 non-null object
host_name                             31467 non-null object
host_since                            31467 non-null object
host_location                         31389 non-null object
host_about                            19477 non-null object
host_response_time                    23183 non-null object
host_response_rate                    23183 non-null object
host_acceptance_rate                  24923 non-null object
host_is_superhost                     31467 non-null object
host_thumbnail_url                    31467 non-null object
host_picture_url                      31467 non-null object
host_neighbourhood                    25331 non-null object
host_listings_count                   31467 non-null float64
host_total_listings_count             31467 non-null float64
host_verifications                    31471 non-null object
host_has_profile_pic                  31467 non-null object
host_identity_verified                 31467 non-null object
neighbourhood                         20324 non-null object
neighbourhood_cleansed                 31471 non-null object
neighbourhood_group_cleansed           31471 non-null object
latitude                              31471 non-null float64
longitude                             31471 non-null float64
property_type                         31471 non-null object
room_type                             31471 non-null object
accommodates                          31471 non-null int64
bathrooms                             0 non-null float64
bathrooms_text                        31440 non-null object
bedrooms                              27831 non-null float64
beds                                  31162 non-null float64
amenities                             31471 non-null object
price                                 31471 non-null object
minimum_nights                        31471 non-null int64
maximum_nights                        31471 non-null int64
minimum_minimum_nights                 31469 non-null float64
maximum_minimum_nights                 31469 non-null float64
minimum_maximum_nights                 31469 non-null float64
maximum_maximum_nights                 31469 non-null float64
minimum_nights_avg_ntm                 31469 non-null float64
maximum_nights_avg_ntm                 31469 non-null float64
calendar_updated                       0 non-null float64
has_availability                       31471 non-null object
availability_30                        31471 non-null int64
availability_60                        31471 non-null int64
availability_90                        31471 non-null int64
availability_365                       31471 non-null int64

```

calendar_last_scraped	31471 non-null object
number_of_reviews	31471 non-null int64
number_of_reviews_ltm	31471 non-null int64
number_of_reviews_l30d	31471 non-null int64
first_review	24244 non-null object
last_review	24244 non-null object
review_scores_rating	23859 non-null float64
review_scores_accuracy	23726 non-null float64
review_scores_cleanliness	23726 non-null float64
review_scores_checkin	23716 non-null float64
review_scores_communication	23724 non-null float64
review_scores_location	23713 non-null float64
review_scores_value	23709 non-null float64
license	6549 non-null object
instant_bookable	31471 non-null object
calculated_host_listings_count	31471 non-null int64
calculated_host_listings_count_entire_homes	31471 non-null int64
calculated_host_listings_count_private_rooms	31471 non-null int64
calculated_host_listings_count_shared_rooms	31471 non-null int64
reviews_per_month	24244 non-null float64

dtypes: float64(22), int64(17), object(35)
memory usage: 17.8+ MB

```
In [5]: df.columns
```

```
Out[5]: Index(['id', 'listing_url', 'scrape_id', 'last_scraped', 'name', 'description',
              'neighborhood_overview', 'picture_url', 'host_id', 'host_url',
              'host_name', 'host_since', 'host_location', 'host_about',
              'host_response_time', 'host_response_rate', 'host_acceptance_rate',
              'host_is_superhost', 'host_thumbnail_url', 'host_picture_url',
              'host_neighbourhood', 'host_listings_count',
              'host_total_listings_count', 'host_verifications',
              'host_has_profile_pic', 'host_identity_verified', 'neighbourhood',
              'neighbourhood_cleansed', 'neighbourhood_group_cleansed', 'latitude',
              'longitude', 'property_type', 'room_type', 'accommodates', 'bathrooms',
              'bathrooms_text', 'bedrooms', 'beds', 'amenities', 'price',
              'minimum_nights', 'maximum_nights', 'minimum_minimum_nights',
              'maximum_minimum_nights', 'minimum_maximum_nights',
              'maximum_maximum_nights', 'minimum_nights_avg_ntm',
              'maximum_nights_avg_ntm', 'calendar_updated', 'has_availability',
              'availability_30', 'availability_60', 'availability_90',
              'availability_365', 'calendar_last_scraped', 'number_of_reviews',
              'number_of_reviews_ltm', 'number_of_reviews_l30d', 'first_review',
              'last_review', 'review_scores_rating', 'review_scores_accuracy',
              'review_scores_cleanliness', 'review_scores_checkin',
              'review_scores_communication', 'review_scores_location',
              'review_scores_value', 'license', 'instant_bookable',
              'calculated_host_listings_count',
              'calculated_host_listings_count_entire_homes',
              'calculated_host_listings_count_private_rooms',
              'calculated_host_listings_count_shared_rooms', 'reviews_per_month'],
              dtype='object')
```

```
In [6]: #Examining columns that can be converted to dummy variables
df['host_response_time'].unique()
```

```
Out[6]: array([nan, 'within an hour', 'within a few hours', 'within a day',
               'a few days or more'], dtype=object)
```

I still want to add host_response and host_acceptance rates even though they are object type. They seem to have relevance. Start by stripping "%" and converting each value to a float.

```
In [12]: df['host_response_rate'] = df['host_response_rate'].str.rstrip('%').astype(
         'float') / 100.0
         df['host_acceptance_rate'] = df['host_acceptance_rate'].str.rstrip('%').
         astype('float') / 100.0
```

```
In [18]: df['price'] = df['price'].str.lstrip('$')
```

```
In [20]: df['price'] = pd.to_numeric(df['price'], errors='coerce')
```

Since I will be performing classification based modeling, I want all the variables to be numbers. I decided to create a dataframe with all existing numericals and then adding categoricals that will later be converted to dummy variables.

```
In [22]: #Select all int and float types
df1 = pd.DataFrame()
df1 = df.select_dtypes(include=['int64', 'float64'])
```

```
In [23]: df1.head(2)
```

Out[23]:

	id	scrape_id	host_id	host_response_rate	host_acceptance_rate	host_listings_count
0	109	20201103044428	521	NaN	0.0	1.0
1	2708	20201103044428	3008	1.0	1.0	2.0

```
In [24]: #Add relevant categoricals/object types that will be converted to dummies
         with get_dummies
df1['host_response_time'] = df['host_response_time']
df1['room_type'] = df['room_type']
```

This is to convert all boolean values to binary encoding.

```
In [25]: #Convert t/f to binary encoding.
df1['instant_bookable'] = df['instant_bookable'].map({'f': 0, 't': 1})
df1['has_availability'] = df['has_availability'].map({'f': 0, 't': 1})
df1['host_identity_verified'] = df['host_identity_verified'].map({'f': 0, 't': 1})
df1['host_has_profile_pic'] = df['host_has_profile_pic'].map({'f': 0, 't': 1})
df1['host_is_superhost'] = df['host_is_superhost'].map({'f': 0, 't': 1})
```

```
In [26]: df1.head(2)
```

Out[26]:

	id	scrape_id	host_id	host_response_rate	host_acceptance_rate	host_listings_count
0	109	20201103044428	521	NaN	0.0	1.0
1	2708	20201103044428	3008	1.0	1.0	2.0

```
In [27]: #One hot encode categoricals (neighbourhood_group_cleansed, host_respons  
e_time), drop_first = True for logistic  
#Although not required for the other models I will be running  
df2 = pd.get_dummies(df1)
```

Null values

Null values for this data set were particularly interesting since some of the columns didn't have a clear value to replace the NaN values with. If I were to simply replace NaN's with any measure of central tendency, the distribution of that column/feature would skew heavily. Also removing the NaN's would potentially remove a lot of important information as many columns had around 8,000/31,000 null values.

A count of the null values will be displayed below. However, as there were also many columns with few null values (2, 4) I made the decision to just drop those rows since there wouldn't be much impact on the distribution of the data for this case. 2/31,000.

```
In [28]: #Clearly there are a lot of null values that need to be dealt with.  
df2.head(2)
```

Out[28]:

	id	scrape_id	host_id	host_response_rate	host_acceptance_rate	host_listings_count
0	109	20201103044428	521	NaN	0.0	1.0
1	2708	20201103044428	3008	1.0	1.0	2.0


```
In [29]: df2.isna().sum()
```

```

Out[29]: id 0
        scrape_id 0
        host_id 0
        host_response_rate 8288
        host_acceptance_rate 6548
        host_listings_count 4
        host_total_listings_count 4
        latitude 0
        longitude 0
        accommodates 0
        bathrooms 31471
        bedrooms 3640
        beds 309
        price 889
        minimum_nights 0
        maximum_nights 0
        minimum_minimum_nights 2
        maximum_minimum_nights 2
        minimum_maximum_nights 2
        maximum_maximum_nights 2
        minimum_nights_avg_ntm 2
        maximum_nights_avg_ntm 2
        calendar_updated 31471
        availability_30 0
        availability_60 0
        availability_90 0
        availability_365 0
        number_of_reviews 0
        number_of_reviews_ltm 0
        number_of_reviews_l30d 0
        review_scores_rating 7612
        review_scores_accuracy 7745
        review_scores_cleanliness 7745
        review_scores_checkin 7755
        review_scores_communication 7747
        review_scores_location 7758
        review_scores_value 7762
        calculated_host_listings_count 0
        calculated_host_listings_count_entire_homes 0
        calculated_host_listings_count_private_rooms 0
        calculated_host_listings_count_shared_rooms 0
        reviews_per_month 7227
        instant_bookable 0
        has_availability 0
        host_identity_verified 4
        host_has_profile_pic 4
        host_is_superhost 4
        host_response_time_a few days or more 0
        host_response_time_within a day 0
        host_response_time_within a few hours 0
        host_response_time_within an hour 0
        room_type_Entire home/apt 0
        room_type_Hotel room 0
        room_type_Private room 0
        room_type_Shared room 0
        dtype: int64

```

```
In [30]: df2 = df2.drop(['bathrooms', 'calendar_updated', 'id', 'scrape_id', 'host_id'], axis=1)
```

```
In [34]: #Based on median: 1.0
df2['host_response_rate'].replace(np.NaN, 1, inplace=True)
df2['bedrooms'].replace(np.NaN, 1, inplace=True)
df2['beds'].replace(np.NaN, 1, inplace=True)
df2['review_scores_checkin'].replace(np.NaN, 10, inplace=True)
df2['review_scores_communication'].replace(np.NaN, 10, inplace=True)
df2['review_scores_location'].replace(np.NaN, 10, inplace=True)
df2['review_scores_accuracy'].replace(np.NaN, 10, inplace=True)
```

Host Acceptance rate is interesting it has ~8000/31000 null values so making them a single value would skew the data heavily. The values don't lean toward a specific value so I concluded the best way would be to use the `np.random.choice()` method.

```
In [35]: #To show that values are being evenly distributed based on previous percentage distribution
x = df2['host_acceptance_rate'].value_counts(normalize=True)
print(x)
```

```
1.00    0.274887
0.99    0.088553
0.98    0.067167
0.97    0.044979
0.00    0.038077
...
0.15    0.000120
0.12    0.000120
0.24    0.000080
0.05    0.000040
0.02    0.000040
Name: host_acceptance_rate, Length: 99, dtype: float64
```

```
In [36]: missing = df2['host_acceptance_rate'].isnull()
df2.loc[missing, 'host_acceptance_rate'] = np.random.choice(x.index, size=len(df2[missing]), p=x.values)
```

```
In [37]: df2['host_acceptance_rate'].value_counts(normalize=True)
```

```
Out[37]: 1.00    0.274221
0.99    0.089130
0.98    0.067205
0.97    0.045502
0.00    0.038162
...
0.07    0.000095
0.12    0.000095
0.24    0.000095
0.05    0.000032
0.02    0.000032
Name: host_acceptance_rate, Length: 99, dtype: float64
```

The following is repetitive information that fills in ~8,000 null values for each of these specified columns using random choice to evenly replace NaNs based on percentages each unique values appearance (Hence, normalize=True).

```
In [38]: a = df2['review_scores_rating'].value_counts(normalize=True)
```

```
In [39]: missing1 = df2['review_scores_rating'].isnull()  
df2.loc[missing1, 'review_scores_rating'] = np.random.choice(a.index, size=len(df2[missing1]), p=a.values)
```

```
In [40]: b = df2['review_scores_cleanliness'].value_counts(normalize=True)
```

```
In [41]: missing2 = df2['review_scores_cleanliness'].isnull()  
df2.loc[missing2, 'review_scores_cleanliness'] = np.random.choice(b.index, size=len(df2[missing2]), p=b.values)
```

```
In [42]: c = df2['review_scores_value'].value_counts(normalize=True)
```

```
In [43]: missing3 = df2['review_scores_value'].isnull()  
df2.loc[missing3, 'review_scores_value'] = np.random.choice(c.index, size=len(df2[missing3]), p=c.values)
```

```
In [44]: d = df2['reviews_per_month'].value_counts(normalize=True)
```

```
In [45]: missing4 = df2['reviews_per_month'].isnull()  
df2.loc[missing4, 'reviews_per_month'] = np.random.choice(d.index, size=len(df2[missing4]), p=d.values)
```

```
In [47]: e = df2['price'].value_counts(normalize=True)
```

```
In [48]: missing5 = df2['price'].isnull()  
df2.loc[missing5, 'price'] = np.random.choice(e.index, size=len(df2[missing5]), p=e.values)
```

Moving forward, with the final few null values (2, 4) in certain columns. I will be just dropping the rows that contain null values as it will not have a significant impact on the data.

```
In [49]: #Simple line that drops all remaining null values.  
df2.dropna(inplace=True)
```

```
In [50]: df2.head(2)
```

Out[50]:

	host_response_rate	host_acceptance_rate	host_listings_count	host_total_listings_count	latitude
0	1.0	0.0	1.0	1.0	33.9821
1	1.0	1.0	2.0	2.0	34.0971

```
In [51]: df2.isna().sum()
```

```
Out[51]: host_response_rate      0
host_acceptance_rate            0
host_listings_count             0
host_total_listings_count       0
latitude                       0
longitude                      0
accommodates                    0
bedrooms                       0
beds                           0
price                          0
minimum_nights                  0
maximum_nights                  0
minimum_minimum_nights         0
maximum_minimum_nights         0
minimum_maximum_nights         0
maximum_maximum_nights         0
minimum_nights_avg_ntm         0
maximum_nights_avg_ntm         0
availability_30                 0
availability_60                 0
availability_90                 0
availability_365                0
number_of_reviews               0
number_of_reviews_ltm           0
number_of_reviews_l30d          0
review_scores_rating             0
review_scores_accuracy           0
review_scores_cleanliness        0
review_scores_checkin            0
review_scores_communication      0
review_scores_location           0
review_scores_value              0
calculated_host_listings_count  0
calculated_host_listings_count_entire_homes 0
calculated_host_listings_count_private_rooms 0
calculated_host_listings_count_shared_rooms 0
reviews_per_month                0
instant_bookable                 0
has_availability                 0
host_identity_verified           0
host_has_profile_pic             0
host_is_superhost                0
host_response_time_a few days or more 0
host_response_time_within a day      0
host_response_time_within a few hours 0
host_response_time_within an hour    0
room_type_Entire home/apt          0
room_type_Hotel room               0
room_type_Private room             0
room_type_Shared room              0
dtype: int64
```

```
In [52]: df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 31465 entries, 0 to 31470
Data columns (total 50 columns):
host_response_rate      31465 non-null float64
host_acceptance_rate    31465 non-null float64
host_listings_count      31465 non-null float64
host_total_listings_count 31465 non-null float64
latitude                 31465 non-null float64
longitude                31465 non-null float64
accommodates             31465 non-null int64
bedrooms                 31465 non-null float64
beds                     31465 non-null float64
price                    31465 non-null float64
minimum_nights            31465 non-null int64
maximum_nights            31465 non-null int64
minimum_minimum_nights    31465 non-null float64
maximum_minimum_nights    31465 non-null float64
minimum_maximum_nights    31465 non-null float64
maximum_maximum_nights    31465 non-null float64
minimum_nights_avg_ntm    31465 non-null float64
maximum_nights_avg_ntm    31465 non-null float64
availability_30           31465 non-null int64
availability_60           31465 non-null int64
availability_90           31465 non-null int64
availability_365          31465 non-null int64
number_of_reviews         31465 non-null int64
number_of_reviews_ltm     31465 non-null int64
number_of_reviews_l30d    31465 non-null int64
review_scores_rating       31465 non-null float64
review_scores_accuracy     31465 non-null float64
review_scores_cleanliness  31465 non-null float64
review_scores_checkin      31465 non-null float64
review_scores_communication 31465 non-null float64
review_scores_location     31465 non-null float64
review_scores_value        31465 non-null float64
calculated_host_listings_count 31465 non-null int64
calculated_host_listings_count_entire_homes 31465 non-null int64
calculated_host_listings_count_private_rooms 31465 non-null int64
calculated_host_listings_count_shared_rooms 31465 non-null int64
reviews_per_month         31465 non-null float64
instant_bookable          31465 non-null int64
has_availability           31465 non-null int64
host_identity_verified     31465 non-null float64
host_has_profile_pic       31465 non-null float64
host_is_superhost          31465 non-null float64
host_response_time_a few days or more 31465 non-null uint8
host_response_time_within a day 31465 non-null uint8
host_response_time_within a few hours 31465 non-null uint8
host_response_time_within an hour 31465 non-null uint8
room_type_Entire home/apt  31465 non-null uint8
room_type_Hotel room       31465 non-null uint8
room_type_Private room     31465 non-null uint8
room_type_Shared room      31465 non-null uint8
dtypes: float64(26), int64(16), uint8(8)
memory usage: 10.6 MB

```


Here I will be classifying each of our classes as a number. There are three classes that hint to ternary-class models: and these classes are binned based on rating performance. I have labeled the classes 0 (subpar), 1 (great), 2(best) based on the ranges (0-94], (94-99], (99-100], respectively.

The process I took when experimenting with unique ranges.

Old:

```
df_best = df.loc[(df['review_scores_rating'] == 100)]
```

```
df_great = df.loc[(df['review_scores_rating'] < 100) &  
(df['review_scores_rating'] >= 95)]
```

```
df_good = df.loc[(df['review_scores_rating'] >= 80) &  
(df['review_scores_rating'] < 95)]
```

```
df_bad = df.loc[(df['review_scores_rating'] < 80)]
```

New:

```
df_best = df.loc[(df['review_scores_rating'] == 100)]
```

```
df_great = df.loc[(df['review_scores_rating'] < 100) &  
(df['review_scores_rating'] >= 95)]
```

```
df_subpar = df.loc[(df['review_scores_rating'] < 95)]
```

Descriptions:

Best: The best score of 100

Great: 95-100, contains the median, mode

Supbar: Lower than the mean of 95

```
In [53]: df2['target'] = pd.cut(df2['review_scores_rating'], [0, 95, 99, 100], right=True, \
                                labels=['0', '1', '2'])
```

```
In [54]: df2.drop('review_scores_rating', axis=1, inplace=True)
```

```
In [55]: df2['target'].value_counts()
```

```
Out[55]: 0    11662
         1    10479
         2     9324
         Name: target, dtype: int64
```

As you can see, my cutoffs for each class are not only suitable ranges, but also address class imbalance. I made these cutoff decisions based on measures of central tendency where the 'best'(2) class has only perfect scores, the 'great'(1) class has scores greater than the mean up to, but not including, the best (median and mode) and the 'subpar'(0) class contains all values lower than the mean.


Some additional explanation: After conducting side research on airbnb ratings, all values lower than the mean potentially represent ratings where customers had felt their experience was not the best and could have complaints/suggestions to make.

Class imbalance: For future datasets, my solution to class imbalance will not particularly be the same as it is here. I would definitely implement SMOTE (which creates artificial data in underrepresented classes).

```
In [56]: #Binned using .cut; however, still category type. Converting to numeric
df2['target'] = pd.to_numeric(df2['target'], errors='coerce')
```

```
In [57]: #Visualization of target class distribution.
df2['target'].plot.hist(bins=5)
```

```
Out[57]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1e7e71d0>
```



```
In [58]: df2.head(2)
```

```
Out[58]:
```

	host_response_rate	host_acceptance_rate	host_listings_count	host_total_listings_count	latitude
0	1.0	0.0	1.0	1.0	33.9821
1	1.0	1.0	2.0	2.0	34.0971




```
In [59]: #Dropping unnecessary or redundant/repetitive columns
df2.drop(['host_listings_count', 'host_total_listings_count', 'minimum_minimum_nights', 'maximum_minimum_nights', \
          'minimum_maximum_nights', 'maximum_maximum_nights', 'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', \
          'number_of_reviews_ltm', 'number_of_reviews_l30d', 'calculated_host_listings_count', \
          'calculated_host_listings_count_entire_homes', 'calculated_host_listings_count_private_rooms', \
          'calculated_host_listings_count_shared_rooms'], axis=1, inplace=True)
```




```
In [60]: #Outlier features checks: ACCOMODATES, BEDROOMS, BEDS, MINIMUM_NIGHTS(remove single values)
df2['beds'].value_counts()
```


```
Out[60]: 1.0    15800
2.0     7160
3.0     3191
4.0     2121
0.0     1210
5.0      895
6.0      561
7.0      211
8.0      144
9.0       66
10.0      42
12.0      16
11.0      16
14.0      11
16.0       4
22.0       3
13.0       3
19.0       2
15.0       2
32.0       1
50.0       1
44.0       1
18.0       1
20.0       1
21.0       1
23.0       1
Name: beds, dtype: int64
```

```
In [61]: #Create a dataframe with non binary/discrete values to look at boxplot distribution and remove any outliers.  
non_bin = df2[['accommodates', 'bedrooms', 'beds', 'minimum_nights', 'maximum_nights', 'number_of_reviews', \  
               'reviews_per_month']]
```




```
In [62]: for column in non_bin:
          non_bin[column].plot.hist(density=True, label= column+' histogram')
          non_bin[column].plot.kde(label=column+' kde')
          plt.legend()
          plt.show()
```

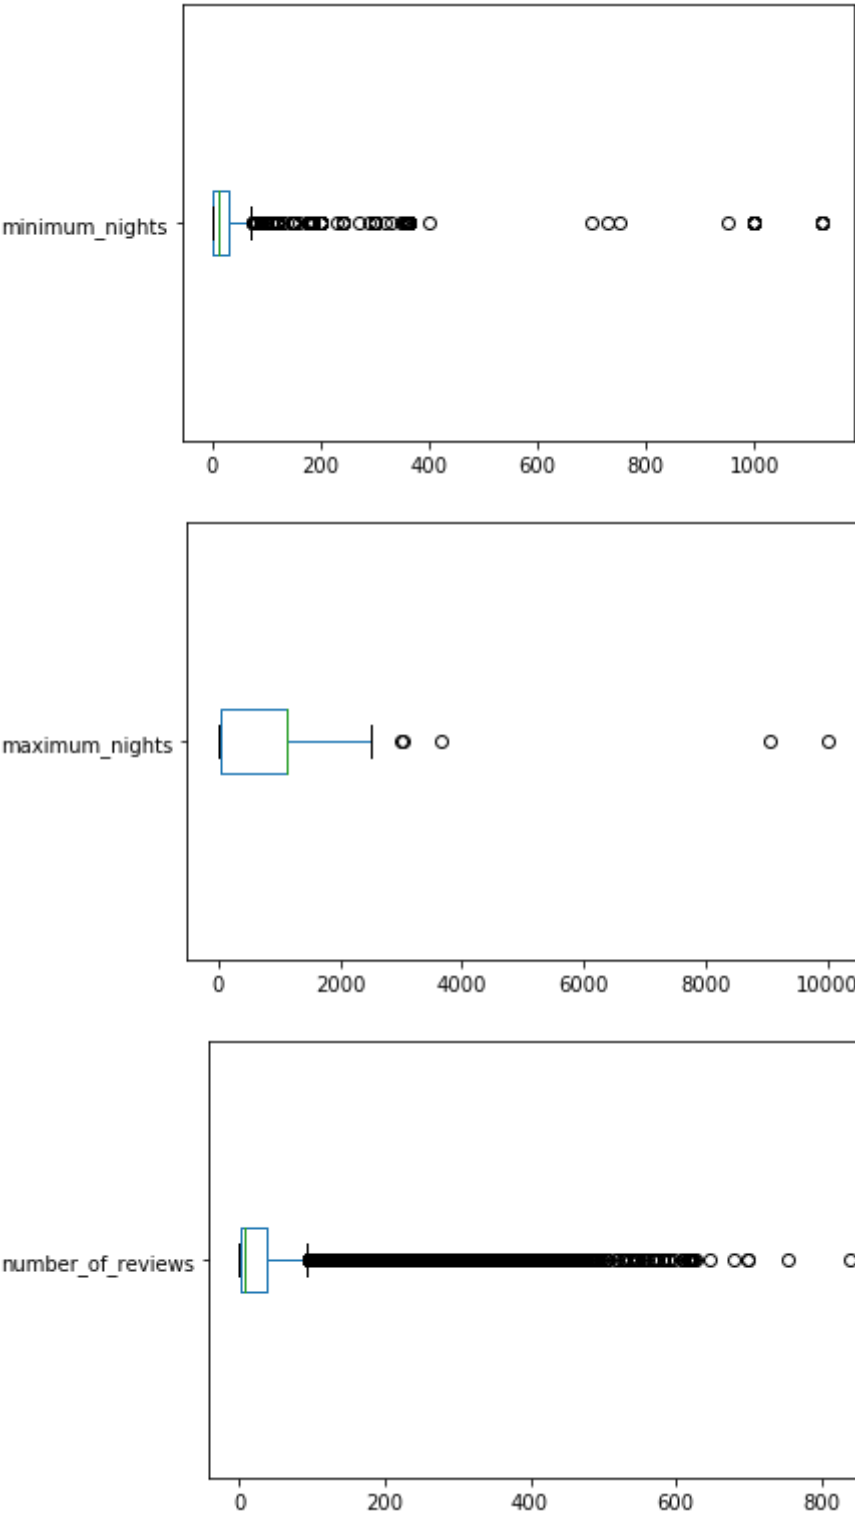








```
In [63]: for column in non_bin:
          non_bin[column].plot.box(vert=False)
          plt.show()
```









```
In [64]: sns.boxplot(x=df2['beds'])
```

```
Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2e2c25f8>
```



```
In [65]: sns.boxplot(x=df2['price'])
```

```
Out[65]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2fcd7a58>
```



Removing outliers from each column based on visual boxplots:

```
In [66]: df2 = df2.loc[(df2['beds'] < 10)]
```

```
In [67]: df2 = df2.loc[(df2['minimum_nights'] < 80)]
```

```
In [68]: df2 = df2.loc[(df2['accommodates'] < 12)]
```

```
In [69]: df2 = df2.loc[(df2['bedrooms'] < 10)]
```

```
In [70]: df2 = df2.loc[(df2['maximum_nights'] < 5000)]
```

```
In [71]: df2 = df2.loc[(df2['number_of_reviews'] < 600)]
```

```
In [72]: df2 = df2.loc[(df2['reviews_per_month'] < 15)]
```

```
In [73]: df2 = df2.loc[(df2['price'] < 600)]
```

EDA


Exploratory Data Analysis (EDA) here consists of creating visualizations in order to understand and summarize key characteristics of my data. I will mostly be illustrating visuals on what my features/columns look across my three classes. Specifically, looking for variance/seperability within classes so the model can make better predictions and can theoretically distinguish the classes better.

```
In [75]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29771 entries, 0 to 31470
Data columns (total 36 columns):
host_response_rate                29771 non-null float64
host_acceptance_rate              29771 non-null float64
latitude                          29771 non-null float64
longitude                         29771 non-null float64
accommodates                      29771 non-null int64
bedrooms                         29771 non-null float64
beds                             29771 non-null float64
price                             29771 non-null float64
minimum_nights                   29771 non-null int64
maximum_nights                   29771 non-null int64
availability_30                   29771 non-null int64
availability_60                   29771 non-null int64
availability_90                   29771 non-null int64
availability_365                  29771 non-null int64
number_of_reviews                 29771 non-null int64
review_scores_accuracy            29771 non-null float64
review_scores_cleanliness         29771 non-null float64
review_scores_checkin             29771 non-null float64
review_scores_communication       29771 non-null float64
review_scores_location            29771 non-null float64
review_scores_value               29771 non-null float64
reviews_per_month                 29771 non-null float64
instant_bookable                  29771 non-null int64
has_availability                  29771 non-null int64
host_identity_verified            29771 non-null float64
host_has_profile_pic              29771 non-null float64
host_is_superhost                 29771 non-null float64
host_response_time_a few days or more 29771 non-null uint8
host_response_time_within a day      29771 non-null uint8
host_response_time_within a few hours 29771 non-null uint8
host_response_time_within an hour     29771 non-null uint8
room_type_Entire home/apt          29771 non-null uint8
room_type_Hotel room               29771 non-null uint8
room_type_Private room             29771 non-null uint8
room_type_Shared room              29771 non-null uint8
target                             29771 non-null int64
dtypes: float64(17), int64(11), uint8(8)
memory usage: 6.8 MB
```


```
In [76]: #The perfect class receives marginally more reviews than the other 2 classes.  
pd.crosstab(df2['reviews_per_month'], df2['target']).plot.bar(figsize=(10,10))
```

Out[76]: <matplotlib.axes._subplots.AxesSubplot at 0x1a301be4e0>




```
In [54]: pd.crosstab(df2['review_scores_communication'], df2['target']).plot.bar(
figsize=(10,10))
#Noteable trend, for the subpar class, although it has some perfect scores,
it has scores for every value down to 2 more
#prominent than the other, better classes.
#So, communication was not that great.
```

Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x1a22905ef0>





```
In [55]: pd.crosstab(df2['review_scores_cleanliness'], df2['target']).plot.bar(figsize=(10,10))  
#Similar trend, 1,2 have a lot more perfect scores than 0 and 0 is more prominent in the other scores down to 2.
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1d9e2898>
```




```
In [56]: pd.crosstab(df2['review_scores_accuracy'], df2['target']).plot.bar(figsize=(10,10))
```

```
Out[56]: <matplotlib.axes._subplots.AxesSubplot at 0x1a24122358>
```




```
In [57]: pd.crosstab(df2['review_scores_location'], df2['target']).plot.bar(figsize=(10,10))
```

```
Out[57]: <matplotlib.axes._subplots.AxesSubplot at 0x1a23be9240>
```




```
In [58]: pd.crosstab(df2['review_scores_checkin'], df2['target']).plot.bar(figsize=(10,10))
```

```
Out[58]: <matplotlib.axes._subplots.AxesSubplot at 0x1a25bca7f0>
```




```
In [59]: pd.crosstab(df2['availability_30'], df2['target']).plot.bar(figsize=(10, 10))  
#As you can see within 0 days, essentially none of the classes are booked. But, as you progress up to 30, class 0 is  
#more available and we can conclude that less people book the lower rate d listings.
```

Out[59]: <matplotlib.axes._subplots.AxesSubplot at 0x1a24bf3b38>




```
In [60]: pd.crosstab(df2['availability_60'], df2['target']).plot.bar(figsize=(10, 10))
#Very similar trend to availability within 30 days
```

```
Out[60]: <matplotlib.axes._subplots.AxesSubplot at 0x1a24ba2b38>
```




```
In [61]: pd.crosstab(df2['has_availability'], df2['target']).plot.bar(figsize=(10,10))  
#This graph tells us the lower rating listings have more availability versus the better rating listings.  
#Which intuitively makes perfect sense. The better listings will have lower availability since they are being booked.
```

Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2ac725c0>




```
In [62]: pd.crosstab(df2['instant_bookable'], df2['target']).plot.bar(figsize=(10,10))  
#The better rated listings are probably less "instantly bookable" because  
#people plan ahead to reserve better listings
```

Out[62]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2bcda240>





```
In [63]: pd.crosstab(df2['accommodates'], df2['target']).plot.bar(figsize=(10,10))
```

```
Out[63]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2bd12dd8>
```




```
In [64]: pd.crosstab(df2['host_is_superhost'], df2['target']).plot.bar(figsize=(10,10))  
#Classes 1 and 2 have way more superhost status than class 0. A superhost is a verified host that provides excellent  
#experiences and stays and is a "role model".
```

Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2bd32d30>



```
In [65]: pd.crosstab(df2['host_identity_verified'], df2['target']).plot.bar(figsize=(10,10))
```


```
Out[65]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2bd32a90>
```



Exploring different types of graphs


```
In [66]: sns.displot(data=df2, x='availability_30', hue='target', kind='kde')
```

```
Out[66]: <seaborn.axisgrid.FacetGrid at 0x1a2c098e10>
```




```
In [67]: sns.displot(data=df2, x='availability_60', hue='target', kind='kde')
```

```
Out[67]: <seaborn.axisgrid.FacetGrid at 0x1a178b24e0>
```




```
In [68]: sns.displot(data=df2, x='availability_90', hue='target', kind='kde')
```

```
Out[68]: <seaborn.axisgrid.FacetGrid at 0x1a178b2470>
```



```
In [69]: sns.displot(data=df2, x='availability_365', hue='target', kind='kde')
```

```
Out[69]: <seaborn.axisgrid.FacetGrid at 0x1a2f9094a8>
```




This graph helps visualize the availability within 30, 60, 90, 365 days better than the barplots do. A trend to notice is that class 0 (subpar) consistently has more availability while class 2 (best) is always booked and has the least availability out of the classes.

```
In [70]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30449 entries, 0 to 31470
Data columns (total 35 columns):
host_response_rate          30449 non-null float64
host_acceptance_rate        30449 non-null float64
latitude                    30449 non-null float64
longitude                   30449 non-null float64
accommodates                30449 non-null int64
bedrooms                   30449 non-null float64
beds                       30449 non-null float64
minimum_nights              30449 non-null int64
maximum_nights              30449 non-null int64
availability_30              30449 non-null int64
availability_60              30449 non-null int64
availability_90              30449 non-null int64
availability_365             30449 non-null int64
number_of_reviews            30449 non-null int64
review_scores_accuracy        30449 non-null float64
review_scores_cleanliness     30449 non-null float64
review_scores_checkin         30449 non-null float64
review_scores_communication   30449 non-null float64
review_scores_location        30449 non-null float64
review_scores_value           30449 non-null float64
reviews_per_month             30449 non-null float64
instant_bookable             30449 non-null int64
has_availability              30449 non-null int64
host_identity_verified        30449 non-null float64
host_has_profile_pic          30449 non-null float64
host_is_superhost             30449 non-null float64
host_response_time_a few days or more 30449 non-null uint8
host_response_time_within a day      30449 non-null uint8
host_response_time_within a few hours 30449 non-null uint8
host_response_time_within an hour    30449 non-null uint8
room_type_Entire home/apt         30449 non-null uint8
room_type_Hotel room              30449 non-null uint8
room_type_Private room           30449 non-null uint8
room_type_Shared room            30449 non-null uint8
target                          30449 non-null int64
dtypes: float64(16), int64(11), uint8(8)
memory usage: 6.7 MB
```


```
In [71]: sns.displot(data=df2, x='maximum_nights', hue='target', kind='kde')
```

```
Out[71]: <seaborn.axisgrid.FacetGrid at 0x1a2f8f0c88>
```




```
In [72]: sns.displot(data=df2, x='minimum_nights', hue='target', kind='kde')
```

```
Out[72]: <seaborn.axisgrid.FacetGrid at 0x1a2318a2e8>
```




```
In [73]: sns.displot(data=df2, x='host_acceptance_rate', hue='target', kind='kde')
#Maybe the lower rating listings accept for people because they need money? And therefore the quality of the listing
#is lower. Conversely, the best rating listings are more strict on the rating of the guest and whether or not they
#are allowed to stay.
```

Out[73]: <seaborn.axisgrid.FacetGrid at 0x1a231ef6d8>





```
In [74]: sns.displot(data=df2, x='host_response_rate', hue='target', kind='kde')
```

```
Out[74]: <seaborn.axisgrid.FacetGrid at 0x1a2a172860>
```




```
In [75]: sns.displot(data=df2, x='number_of_reviews', hue='target', kind='kde')  
#Class 2 seems to have a lot of listings with ~0-30 reviews by a large margin.
```

```
Out[75]: <seaborn.axisgrid.FacetGrid at 0x1a2a195b38>
```



```
In [78]: sns.displot(data=df2, x='price', hue='target', kind='kde')
#While the classes seem to remain consistent as price increases, we CAN
#say that class 0 has more listings in the
#less expensive range
```

Out[78]: <seaborn.axisgrid.FacetGrid at 0x1a206d3940>



```
In [79]: df2.head(2)
```

Out[79]:

	host_response_rate	host_acceptance_rate	latitude	longitude	accommodates	bedrooms	bi
0	1.0	0.0	33.98209	-118.38494	6	2.0	
1	1.0	1.0	34.09768	-118.34602	1	1.0	

```
In [ ]:
```

Initial Modeling

```
In [80]: #Define your X and Y for train, test split
X = df2.drop(['target'],axis=1)
y = df2['target']
```

```
In [81]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

KNN

KNN (K-Nearest Neighbors) modeling. I implemented a function that will help determine the best K value based on accuracy of the models.

```
In [82]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [83]: def find_best_k(X_train, y_train, X_test, y_test, min_k=1, max_k=25):
    best_k = 0
    best_score = 0.0
    for k in range(min_k, max_k+1, 2):
        knn = KNeighborsClassifier(n_neighbors=k)
        knn.fit(X_train, y_train)
        preds = knn.predict(X_test)
        acc = accuracy_score(y_test, preds)
        if acc > best_score:
            best_k = k
            best_score = acc

    print("Best Value for k: {}".format(best_k))
    print("acc: {}".format(best_score))
```

```
In [84]: #This is for KNN because it uses distance so if data is not on the same
scale, it will be weighted differently
from sklearn.preprocessing import StandardScaler
```

```
In [85]: #Transforming X_train and X_test using StandardScaler
#Fit is where StandardScaler understand the distribution and variance of
the data a learns how to scale it.
#Transform then applies the "memorized" fit and applies the scaling.
scaler = StandardScaler()
scaler.fit(X_train)


X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

#Running a baseline KNN model
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

y_predict = knn.predict(X_test)
```


```
In [86]: plot_confusion_matrix(knn, X_train, y_train)
```

```
Out[86]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a23c975f8>
```



```
In [87]: plot_confusion_matrix(knn, X_test, y_test)
```

```
Out[87]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2b2c8ac8>
```



```
In [88]: print(classification_report(y_train, knn.predict(X_train)))
print(classification_report(y_test, knn.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.73	0.78	0.76	8321
1	0.67	0.70	0.69	7445
2	0.69	0.59	0.64	6562
accuracy			0.70	22328
macro avg	0.70	0.69	0.69	22328
weighted avg	0.70	0.70	0.70	22328

	precision	recall	f1-score	support
0	0.61	0.64	0.62	2722
1	0.53	0.56	0.54	2504
2	0.48	0.42	0.45	2217
accuracy			0.55	7443
macro avg	0.54	0.54	0.54	7443
weighted avg	0.54	0.55	0.54	7443

```
In [89]: #Testing values for the best k using defined function above.
find_best_k(X_train, y_train, X_test, y_test)
```


Best Value for k: 25
acc: 0.5851135294907968

```
In [90]: #Running adjusted KNN parameters model.
knn = KNeighborsClassifier(n_neighbors=25)
knn.fit(X_train, y_train)

y_predict = knn.predict(X_test)
```


```
In [91]: plot_confusion_matrix(knn, X_train, y_train)
```

```
Out[91]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2c18f240>
```



```
In [92]: plot_confusion_matrix(knn, X_test, y_test)
```

```
Out[92]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2b04f668>
```



```
In [93]: print(classification_report(y_train, knn.predict(X_train)))
print(classification_report(y_test, knn.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.73	0.64	0.68	8321
1	0.61	0.62	0.61	7445
2	0.53	0.60	0.56	6562
accuracy			0.62	22328
macro avg	0.62	0.62	0.62	22328
weighted avg	0.63	0.62	0.62	22328

	precision	recall	f1-score	support
0	0.69	0.60	0.65	2722
1	0.58	0.59	0.59	2504
2	0.49	0.55	0.52	2217
accuracy			0.59	7443
macro avg	0.59	0.58	0.58	7443
weighted avg	0.59	0.59	0.59	7443

```
In [94]: print("Testing Accuracy: {}".format(accuracy_score(y_test, y_predict)))
```

Testing Accuracy: 0.5851135294907968

Improved by 4% in terms of model accuracy. The adjusted model performs better by .05 looking at the weighted average. The F1 score also increased slightly.

In []:

Decision Trees: Final Model

Decision Tree modeling.

```
In [95]: from sklearn.tree import DecisionTreeClassifier
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

```
In [96]: #Instantiate classifier
ctree = DecisionTreeClassifier()

ctree.fit(X_train, y_train)
```


```
Out[96]: DecisionTreeClassifier()
```

```
In [97]: ctree_ypred = ctree.predict(X_test)
```

Running a baseline model with default hyperparameters


```
In [98]: plot_confusion_matrix(ctree, X_train, y_train)
```

```
Out[98]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2b3990f0>
```



```
In [99]: plot_confusion_matrix(ctree, X_test, y_test)
```

```
Out[99]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2e27feb8>
```




```
In [100]: print(classification_report(y_train, ctree.predict(X_train)))
          print(classification_report(y_test, ctree.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8321
1	1.00	1.00	1.00	7445
2	1.00	1.00	1.00	6562
accuracy			1.00	22328
macro avg	1.00	1.00	1.00	22328
weighted avg	1.00	1.00	1.00	22328

	precision	recall	f1-score	support
0	0.64	0.65	0.65	2722
1	0.57	0.57	0.57	2504
2	0.58	0.57	0.57	2217
accuracy			0.60	7443
macro avg	0.60	0.60	0.60	7443
weighted avg	0.60	0.60	0.60	7443


Interestingly, our training set predicted perfectly with 1.0's across all metrics while the testing set specifically had an accuracy of 60%. These indications tell me that there are signs of overfitting, I safely assumed this was the result of some hyperparameters that could be tuned.

```
In [102]: #Further analyzing features importances within the decision tree.
          ctree.feature_importances_
```

```
Out[102]: array([0.01407194, 0.03694513, 0.07398842, 0.06305476, 0.01815772,
                 0.00683402, 0.01595318, 0.06175618, 0.02263324, 0.02941858,
                 0.01388428, 0.01551683, 0.02293637, 0.03675446, 0.19665271,
                 0.13668829, 0.05809635, 0.00428663, 0.00678426, 0.00557085,
                 0.03761697, 0.0680748 , 0.00817935, 0.          , 0.00733673,
                 0.00040626, 0.01383361, 0.00092115, 0.00326617, 0.00335865,
                 0.00610037, 0.0037631 , 0.00065564, 0.00416857, 0.00233446])
```

```
In [103]: def plot_feature_importances(model):
n_features = X_train.shape[1]
plt.figure(figsize=(8,8))
plt.barh(range(n_features), model.feature_importances_, align='center')
plt.yticks(np.arange(n_features), X_train.columns.values)
plt.xlabel('Feature importance')
plt.ylabel('Feature')


plot_feature_importances(ctree)
```



Clear sign of overfitting because there is a max_depth of none on the training set. Test set has an accuracy of .61 For adjusting this model, consider GridSearchCV, and changing the criterion to entropy rather than gini.

```
In [104]: #Utilizing a loop to illustrate the best depth in a cross validation of  
models.  
from sklearn.model_selection import cross_val_score  
score = cross_val_score(ctree, X_train, y_train, cv = 20)  
score.mean()  
depth_range = range(1, 20)  
val = []  
for depth in depth_range:  
    ctree = DecisionTreeClassifier(max_depth = depth)  
    depth_score = cross_val_score(ctree, X_train, y_train, cv = 20)  
    val.append(depth_score.mean())  
print(val)  
plt.figure(figsize=(10,10))  
plt.plot(depth_range, val)  
plt.xlabel('depth range')  
plt.ylabel('CV values')  
plt.show()
```

```
[0.4888035348138736, 0.6236119133110642, 0.6503916340171286, 0.65872203
93206329, 0.667008203296721, 0.6750249885285406, 0.6744886777498612, 0.
6719352752989799, 0.6679934652791817, 0.6648582673122772, 0.66046935114
85898, 0.655811698000597, 0.6490501551454709, 0.6439882734410848, 0.634
1350920765105, 0.6288960044024733, 0.6235218663663231, 0.61792218981334
41, 0.6125485330971656]
```



```
In [105]: from sklearn.model_selection import GridSearchCV
```

```
In [106]: #Using GridSearchCV
tree_params = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, '6'],
    'max_features': ['auto', 'sqrt', 'log2'],
    'min_samples_split': [None, 5, 7, 9, 10]
}
```

```
In [107]: gs_tree = GridSearchCV(ctree, tree_params, cv=3)

gs_tree.fit(X_train, y_train)

print(f"Training Accuracy: {gs_tree.best_score_ :.2%}")
print("")
print(f"Optimal Parameters: {gs_tree.best_params_}")
```

Training Accuracy: 59.92%

Optimal Parameters: {'criterion': 'gini', 'max_depth': None, 'max_features': 'sqrt', 'min_samples_split': 9}

Testing the grid search parameters, there is a training accuracy of 59.92%. While the depth values cross validation gave us 6, I combined some of the grid search values to emphasize a well-rounded model; with hyper parameters from cross validation.


```
In [108]: ctree = DecisionTreeClassifier(max_depth=6, criterion='gini', min_sample
s_split=10)

ctree.fit(X_train, y_train)
```

Out[108]: DecisionTreeClassifier(max_depth=6, min_samples_split=10)


```
In [109]: plot_confusion_matrix(ctree, X_train, y_train)
```

Out[109]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2c957b00>



```
In [110]: plot_confusion_matrix(ctree, X_test, y_test)
```

```
Out[110]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2e16fcf8>
```



```
In [111]: print(classification_report(y_train, ctree.predict(X_train)))
print(classification_report(y_test, ctree.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.70	0.74	0.72	8321
1	0.64	0.71	0.67	7445
2	0.76	0.60	0.67	6562
accuracy			0.69	22328
macro avg	0.70	0.68	0.69	22328
weighted avg	0.69	0.69	0.69	22328

	precision	recall	f1-score	support
0	0.68	0.72	0.70	2722
1	0.63	0.71	0.67	2504
2	0.74	0.59	0.66	2217
accuracy			0.68	7443
macro avg	0.68	0.67	0.67	7443
weighted avg	0.68	0.68	0.68	7443


Accuracy improved by .07 and there is no sign of overfitting as the train and test scores are very similar.

```
In [638]: ctree_ypred = ctree.predict(X_test)
```

So, the final model has an improved accuracy of 68% from 61% with an recall score of .61 for Class 2. The main metric I am looking for in this model is accuracy; however, I thought it would be important to point out that because Class 2 and Class 1 are so close in terms of rating, the ability to distinguish the 2 would be essential.

In terms of misclassified data, with respect to the confusion matrix as well: for Class 0 (2215 predicted correctly and 649 predicted incorrectly). Class 1 (1579 predicted correctly and 938 predicted incorrectly). Class 2 (1370 predicted correctly and 861 predicted incorrectly).

```
In [112]: plot_feature_importances(ctree)
```




In terms of features, clearly the most important feature is `number_of_reviews`. Pre-modeling, I observed that Class 2 which is the perfect class had a marginally greater number of reviews than the other 2 classes. As for host improvements, cleanliness, superhost title and communication seem to stand out.


If we look at percentages, `number_of_reviews`: 40.5%, `accuracy`: 31%, `cleanliness`: 15%, `superhost` title: 3%, `communication`: 1%. While from the baseline, there were lots of contributors. This is because the tuning done on the model directly influenced behavior and minimizing the loss function more efficiently.

```
In [875]: pd.crosstab(df2['review_scores_cleanliness'], df2['target']).plot.bar(figsize=(10,10))
```

```
Out[875]: <matplotlib.axes._subplots.AxesSubplot at 0x1aa7e1cfd0>
```




```
In [637]: #Decision tree visualized.
from sklearn import tree
fig, axes = plt.subplots(nrows = 1,ncols = 1, figsize = (5,5), dpi=300)
tree.plot_tree(ctree,
                feature_names = df2.columns,
                class_names=np.unique(y).astype('str'),
                filled = True)
plt.show()
```



In []:

Random Forests

Random forests modeling.


```
In [113]: from sklearn.ensemble import RandomForestClassifier
```

```
In [114]: rf = RandomForestClassifier()  
  
# Fit classifier  
rf.fit(X_train, y_train)
```

Out[114]: RandomForestClassifier()


```
In [115]: plot_confusion_matrix(rf, X_train, y_train)
```

Out[115]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2c603a20>



```
In [116]: plot_confusion_matrix(rf, X_test, y_test)
```

Out[116]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2c6354e0>



```
In [117]: print(classification_report(y_train, rf.predict(X_train)))
          print(classification_report(y_test, rf.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8321
1	1.00	1.00	1.00	7445
2	1.00	1.00	1.00	6562
accuracy			1.00	22328
macro avg	1.00	1.00	1.00	22328
weighted avg	1.00	1.00	1.00	22328

	precision	recall	f1-score	support
0	0.70	0.71	0.70	2722
1	0.68	0.67	0.67	2504
2	0.66	0.65	0.66	2217
accuracy			0.68	7443
macro avg	0.68	0.68	0.68	7443
weighted avg	0.68	0.68	0.68	7443

```
In [118]: #Using GridSearchCV
          rf_params = {
              'n_estimators': [30, 100],
              'max_depth': [2, 6, 10, 15],
              'min_samples_split': [5, 10],
              'min_samples_leaf': [3, 6]
          }
```

```
In [119]: gs_rf = GridSearchCV(rf, rf_params, cv=3)

          gs_rf.fit(X_train, y_train)

          print(f"Optimal Parameters: {gs_rf.best_params_}")
```

```
Optimal Parameters: {'max_depth': 15, 'min_samples_leaf': 3, 'min_samples_split': 5, 'n_estimators': 100}
```

```
In [715]: from collections import OrderedDict
from sklearn.datasets import make_classification

ensemble_clfs = [
    ("Max_features='sqrt'",
     RandomForestClassifier(warm_start=True, oob_score=True,
                           max_features="sqrt",
                           random_state=123)),
    ("Max_features='log2'",
     RandomForestClassifier(warm_start=True, max_features='log2',
                           oob_score=True,
                           random_state=123)),
    ("Max_features=None",
     RandomForestClassifier(warm_start=True, max_features=None,
                           oob_score=True,
                           random_state=123))
]
```

```

In [716]: error_rate = OrderedDict((label, []) for label, _ in ensemble_clfs)

min_estimators = 10
max_estimators = 200


for label, clf in ensemble_clfs:
    for i in range(min_estimators, max_estimators + 1):
        clf.set_params(n_estimators=i)
        clf.fit(X_train, y_train)

        oob_error = 1 - clf.oob_score_
        error_rate[label].append((i, oob_error))

for label, clf_err in error_rate.items():
    xs, ys = zip(*clf_err)
    plt.plot(xs, ys, label=label)

plt.xlim(min_estimators, max_estimators)
plt.xlabel("n_estimators")
plt.ylabel("OOB error rate")
plt.legend(loc="upper right")
plt.show()

```



```

In [125]: rf = RandomForestClassifier(n_estimators=200, max_depth=15, min_samples_
leaf=3, min_samples_split=10, \
                                     criterion='entropy', max_features='log2')

# Fit classifier
rf.fit(X_train, y_train)

```


```

Out[125]: RandomForestClassifier(criterion='entropy', max_depth=15, max_features
='log2',
                                min_samples_leaf=3, min_samples_split=10,
                                n_estimators=200)

```


```
In [126]: plot_confusion_matrix(rf, X_train, y_train)
```

```
Out[126]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a2e7b56a0>
```



```
In [127]: plot_confusion_matrix(rf, X_test, y_test)
```

```
Out[127]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a5144ff98>
```




```
In [128]: print(classification_report(y_train, rf.predict(X_train)))
print(classification_report(y_test, rf.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.90	0.89	0.89	8321
1	0.85	0.86	0.85	7445
2	0.84	0.84	0.84	6562
accuracy			0.87	22328
macro avg	0.86	0.86	0.86	22328
weighted avg	0.87	0.87	0.87	22328

	precision	recall	f1-score	support
0	0.70	0.72	0.71	2722
1	0.71	0.66	0.68	2504
2	0.64	0.68	0.66	2217
accuracy			0.69	7443
macro avg	0.68	0.68	0.68	7443
weighted avg	0.69	0.69	0.69	7443

```
In [129]: plot_feature_importances(rf)
```



As for random forests, feature importances seem to maintain the same levels as for decision trees. Mainly being number_of_reviews, review_scores_accuracy, cleanliness, superhost title and communication. Of course, the feature importances graph has changed as a result of ensembling multiple decision trees.

XGBoost

XGBoost modeling.

```
In [767]: from xgboost import XGBClassifier
```


```
In [859]: xgb = XGBClassifier()

# Fit classifier
xgb.fit(X_train, y_train)
```

```
Out[859]: XGBClassifier(objective='multi:softprob')
```


```
In [860]: plot_confusion_matrix(xgb, X_train, y_train)
```

```
Out[860]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a742f0320>
```




```
In [861]: plot_confusion_matrix(xgb, X_test, y_test)
```

```
Out[861]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1b4a14fc18>
```



```
In [862]: print(classification_report(y_train, xgb.predict(X_train)))
print(classification_report(y_test, xgb.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.68	0.77	0.72	8489
1	0.68	0.69	0.68	7618
2	0.74	0.60	0.67	6729
accuracy			0.69	22836
macro avg	0.70	0.69	0.69	22836
weighted avg	0.70	0.69	0.69	22836

	precision	recall	f1-score	support
0	0.67	0.75	0.71	2864
1	0.67	0.68	0.67	2517
2	0.74	0.62	0.67	2231
accuracy			0.69	7612
macro avg	0.69	0.68	0.69	7612
weighted avg	0.69	0.69	0.69	7612

```
In [815]: xgb_params = {
    'learning_rate': [0.1, 0.2],
    'max_depth': [6],
    'min_child_weight': [1, 2],
    'subsample': [0.5, 0.7],
    'n_estimators': [100],
}
```

```
In [823]: gs_xgb = GridSearchCV(xgb, xgb_params, scoring='accuracy', cv=None, n_jobs=1)

gs_xgb.fit(X_train, y_train)

print(f"Optimal Parameters: {gs_xgb.best_params_}")
```

Optimal Parameters: {'learning_rate': 0.1, 'max_depth': 6, 'min_child_weight': 1, 'n_estimators': 100, 'subsample': 0.7}


```
In [871]: xgb = XGBClassifier(max_depth=6, n_estimators=200)

# Fit classifier
xgb.fit(X_train, y_train)
```

Out[871]: XGBClassifier(max_depth=6, n_estimators=200, objective='multi:softprob')


```
In [872]: plot_confusion_matrix(xgb, X_train, y_train)
```

Out[872]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1aa792e9b0>



```
In [873]: plot_confusion_matrix(xgb, X_test, y_test)
```

```
Out[873]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a7820cb70>
```



```
In [874]: print(classification_report(y_train, xgb.predict(X_train)))
print(classification_report(y_test, xgb.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.82	0.86	0.84	8489
1	0.79	0.83	0.81	7618
2	0.84	0.75	0.79	6729
accuracy			0.81	22836
macro avg	0.82	0.81	0.81	22836
weighted avg	0.82	0.81	0.81	22836

	precision	recall	f1-score	support
0	0.70	0.73	0.71	2864
1	0.66	0.69	0.68	2517
2	0.72	0.65	0.68	2231
accuracy			0.69	7612
macro avg	0.69	0.69	0.69	7612
weighted avg	0.69	0.69	0.69	7612